

Declarative distributed concurrency in Scala

Sergei Winitzki

Scale by the Bay 2018

November 16, 2018

Talk summary

How I learned to forget semaphores and to love concurrency

Chymyst = an implementation of the Chemical Machine (CM) paradigm

- CM \approx Actors made purely functional and auto-parallelized
- Intuitions about why CM works better than other concurrency models
 - ▶ Comparison with related work: ING Baker, BPMN (workflow)
- New extension for distributed programming: DCM
- Code examples and demos

Not in this talk: academic theory

- Petri nets, π -calculus, join calculus, joinads, mobile agent calculus...
- DCM formulated within some theory of distributed programming

Concurrent & parallel programming: How we cope

Imperative concurrency & parallelism is difficult to reason about:

- low-level API: callbacks, threads, semaphores, mutex locks
- hard to reason about mutable state and running processes
- hard to test – non-deterministic runtime behavior!
 - ▶ race conditions, deadlocks, livelocks

Known declarative approaches to avoid these problems:

Kind of concurrency	Formal structure	Scala implementation
synchronous parallelism	applicative functor	Spark, <code>.par.map()</code>
asynchronous streaming DAG	monadic functor	<code>Future</code> , <code>async/await</code> , RxJava, Akka Streams
unrestricted streaming	recursive monad+	Flink, fs2, ZIO
unrestricted concurrency	?	Akka, Chymyst

For distributed computing: challenges remain

- coordination and consensus, persistence and fault tolerance
- cluster configuration and discovery
 - ▶ distributed coordination as a service: Apache ZooKeeper, etcd